

COMPSCI 589 Homework 2 - Spring 2026

Due **March 12**, 11:55 pm Eastern Time

Instructions

- This homework assignment consists of a programming part.
- While you may discuss problems with your peers (e.g., to discuss high-level approaches), you must answer the questions on your own and implement all solutions independently. In your submission, do explicitly list all students with whom you discussed this assignment.
- Submissions must be typed. You must use $\text{\LaTeX}$ to prepare your submission. Handwritten and scanned submissions will not be accepted.
- The assignment should be submitted on Gradescope as a PDF with marked answers via the Gradescope interface. The source code should be submitted via the Gradescope programming assignment as a .zip file.
- We strongly encourage using Python 3 for your homework code. You may use other languages. In either case, you *must* provide clear instructions for running your code and reproducing your experiments.
- You *may not* use any machine learning-specific libraries in your code, e.g., TensorFlow, PyTorch, or machine learning algorithms implemented in scikit-learn (though you *may* use other functions provided by this library, such as one that splits a dataset into training and testing sets). You may also use libraries such as numpy and matplotlib. If you are not certain whether a specific library is allowed, do ask us.
- All submissions will be checked for plagiarism using two independent plagiarism-detection tools. Renaming variable or function names, moving code within a file, etc., are all strategies that *do not* fool the plagiarism-detection tools we use. Your code will also be compared with submissions from your classmates and with publicly available implementations online.
- If you get caught, all penalties mentioned in the syllabus *will* be applied—which may include directly failing the course with a letter grade of “F”.

→ Before starting this homework, please review this course’s policies on plagiarism by reading the corresponding section of the syllabus.

- The automated system will not accept assignments after 11:55 pm on March 12.
- The TeX file for this homework (which you should use to write your solutions in $\text{\LaTeX}$), as well as the *datasets* you will investigate and the *starting source code* are available on Canvas under “Homework #2 - Supporting Files”. See below for more details.

Programming Assignment [100 Points, total]

In this homework, you will be implementing the Multinomial Naive Bayes algorithm. **Notice that you may not use existing machine learning code for this problem: you must implement the learning algorithms entirely on your own and from scratch.** You will, in particular, train a Naive Bayes algorithm to classify Amazon reviews of CDs and vinyl as either positive or negative. The algorithm will be trained on the Amazon reviews dataset developed by Ni, Li, and McAuley (2019).

Below, we first provide a brief review of the Multinomial Naive Bayes algorithm to help you revisit concepts covered in class, then discuss the structure of the dataset you will explore, and finally outline the experiments and analyses you should conduct.

You can download the dataset and starting source code on Canvas.

1 Multinomial Naive Bayes for Document Classification

A Multinomial Naive Bayes model is trained based on a set of training *documents*, each one belonging to a class. In this assignment, each document corresponds to one product review made by a user and posted on Amazon. You will be given a training set and a test set containing examples of positive product reviews and negative reviews. Recall that in the Multinomial Naive Bayes model, each document is represented by a Bag-of-Words vector $\mathbf{m}$ composed of integer entries; each entry in this vector indicates the frequency of a word in the document (see Fig. 1 for an example).

Assume a classification problem with c classes. Let y_i be the i -th class in the problem, where $1 \leq i \leq c$. Let $\mathbf{Doc}$ be a document, $len(\mathbf{Doc})$ be the number of unique words appearing in $\mathbf{Doc}$, and w_k be the k -th unique word in the document, where $1 \leq k \leq len(\mathbf{Doc})$. Let $\Pr(y_i)$ be the prior probability of class y_i and $\Pr(w_k | y_i)$ be the probability that word w_k appears in documents of class y_i . Recall that when computing the latter probability, the Multinomial Naive Bayes model takes into account the *frequency* with which the word appears in documents of that class, and not just *whether* the word appears in documents of that class (as done, for example, by the Bernoulli Naive Bayes algorithm).

To classify a new document $\mathbf{Doc}$, the algorithm computes, for each class y_i , the following probability:

$$\Pr(y_i | \mathbf{Doc}) = \Pr(y_i) \prod_{k=1}^{len(\mathbf{Doc})} \Pr(w_k | y_i). \quad (1)$$

The algorithm then classifies $\mathbf{Doc}$ as belonging to the class that maximizes the probability in Eq. (1). To train the algorithm, one needs to estimate $\Pr(y_i)$, for each possible class, and $\Pr(w_k | y_i)$, for each possible word and class, given a training set. The prior probability of a class can be estimated, using examples from the training set, as follows:

$$\Pr(y_i) = \frac{N(y_i)}{N}, \quad (2)$$

where N is the total number of documents available in the training set and $N(y_i)$ is the number of documents in the training set that belong to class y_i . Furthermore, the conditional probability of a word given a class can be estimated, using examples from the training set, as follows:

$$\Pr(w_k | y_i) = \frac{n(w_k, y_i)}{\sum_{s=1}^{|V|} n(w_s, y_i)}, \quad (3)$$



figures/HW2_multinomial_repr.png

Figure 1: Example of how a document is represented by the Multinomial Naive Bayes algorithm.

where $n(w_k, y_i)$ is the frequency (total number of occurrences) of word w_k in documents that belong to class y_i , V is the vocabulary containing all unique words that appear in all documents of the training set, and $|V|$ is the length of the vocabulary.

2 The CDs and Vinyl Amazon Reviews Dataset

In this assignment, you will train a machine learning model to classify product reviews as positive or negative. Your algorithm will be trained on an Amazon reviews dataset. This dataset (which is provided to you on Canvas) is split into four files: ‘train-positive’, ‘train-negative’, ‘test-positive’, and ‘test-negative’, containing, respectively, positive and negative instances for training, and positive and negative instances for testing. To train the Multinomial Naive Bayes algorithm, each review/document first needs to be converted to a Bag-of-Words representation, similar to the one shown in Fig. 1. To do so, it is necessary to go through all the examples in the ‘train’ files (including both positive and negative reviews) and construct the vocabulary V of all unique words that appear in the training set. Furthermore, each review that will be analyzed by your

algorithm—either during the training process or when classifying new reviews—needs to be pre-processed: all words in the review should be converted to lowercase, stopwords (i.e., common words such as “the” and “of”) should be removed from reviews, etc. We provide you with a function that performs such pre-processing tasks: `preprocess_text()`, in the `utils.py` file.

Since the dataset to be used in this assignment is relatively large, and you will have to conduct many experiments, we will not ask you to repeatedly manually split the dataset into training and test sets, conduct experiments many times, and evaluate the average performance. Instead, we provide you with one pre-computed split of the dataset; i.e., with one fixed training set and one fixed test set. Each of these sets contains 15,000 examples of positive reviews and 15,000 examples of negative reviews. This means that there is a total of 60,000 product reviews available, 30,000 of which will be used to train the Multinomial Naive Bayes algorithm, and 30,000 of which will be used to test it.

Furthermore, we also provide you with functions to load the product reviews (documents) that will be used as training and test sets: `load_training_set()` and `load_test_set()`, respectively. These functions can be found in the `utils.py` file. Both of them take as input two parameters: the percentage of positive and negative examples that should be randomly drawn from the training or test set, respectively. As an example, if you call `load_training_set(0.5, 0.3)`, the function will return a data structure containing approximately 50% of the existing positive training examples, and approximately 30% of the existing negative training examples. The function `load_test_set()` works similarly: calling `load_test_set(0.2, 0.1)` will return a data structure containing approximately 20% of the existing positive test examples and approximately 10% of the existing negative test examples. This capability is useful for two reasons: *(i)* it allows you to work with very small datasets, at first, while you are still debugging your code, which makes it easier and faster to identify and solve problems; and *(ii)* it will allow you, later, to quantify the impact that unbalanced datasets (i.e., datasets with significantly more examples of one of the classes) have on the performance of the model. You can find an example of how to use the above-mentioned functions in the file `run.py`, included as part of this assignment.

2.1 Example

Consider the following example. This is one of the 15,000 examples of a positive review included as part of the training set. Its contents are:

```
This is a great CD with some old classics. I love listening while
driving down the road. My kids don't hate it either.
```

After pre-processing it (using the `preprocess_text()` function) to remove punctuation, stopwords, etc., it becomes:

```
great cd old classics love listening driving road kids dont hate
either
```

Assume, for simplicity, that the training set contains only two product reviews:

(1)

```
Great box set by Wide Spread Panic. Song selection is good the recording
was done top notch. Very Good and I recommend it highly to those
who are going to buy it.
```

and

(2)

```
I bought both the DVD and the CD. The video is spellbinding and having
the cd in my car is the perfect way to re-live the concert while
I am traveling. I have replayed these so many times that I have
memorized every note. You can't wear this one out!
```

In this case, the data structures returned by the function `load_training_set()` would look something like this:

```
[
  ['great', 'box', 'set', 'wide', 'spread', 'panic', 'song', 'selection',
   'good', 'recording', 'done', 'top', 'knotch', 'good', 'recommend',
   'highly', 'going', 'buy'],

  ['bought', 'dvd', 'cd', 'video', 'spellbinding', 'cd', 'car', 'perfect',
   'way', 'live', 'concert', 'traveling', 'replayed', 'many', 'times',
   'memorized', 'every', 'note', 'cant', 'wear', 'one']
]
```

That is, a vector where each element corresponds to one of the reviews. Each element of the returned vector, in particular, is a Bag-of-Words vector containing the words that appear in the corresponding product review, after all pre-processing steps have been completed.

3 Questions

Hint: While implementing and debugging your solution, we recommend loading only a small portion of the training and test sets. For example, calling the `load_training_set` and `load_test_set` functions with a parameter of 0.0004 will return approximately 0.04% of each dataset—around 10 training and 10 test instances. This will make debugging significantly faster and easier.

- **(Question 1 - 18 Points)** You will first run an experiment to evaluate the performance of the standard Multinomial Naive Bayes algorithm. To do that, your implementation will—as discussed earlier—classify instances by computing posterior probabilities, $\Pr(y_i | Doc)$, according to the standard equation (Eq. (1)). In this experiment, you should use 20% of the training set and 20% of the test set; i.e., call the dataset-loading functions by passing 0.2 as their parameters. First, classify all instances in the test set by comparing the posterior probabilities, $\Pr(y_i | Doc)$, for both classes. Then, report (i) the accuracy of your model; (ii) its precision; (iii) its recall; and (iv) the confusion matrix resulting from this experiment.

Confusion Matrix resulting from the standard Multinomial Naive Bayes algorithm:

True Positive (TP): 2761

False Negative (FN): 254

False Positive (FP): 2360

True Negative (TN): 615

Accuracy of the Model: 0.563606

Precision: 0.539153

Recall: 0.915755

There are 3015 instances of positive reviews. Out of it, 2761 were classified correctly by the standard Multinomial Naive Bayes algorithm

There are 2975 instances of negative reviews. Out of it, 615 were classified correctly by the standard Multinomial Naive Bayes algorithm

56.3606% of the reviews are correctly classified by the standard Multinomial Naive Bayes algorithm

53.9153% of the positive predictions by the standard Multinomial Naive Bayes algorithm are correct

91.5755% of the positive reviews are classified correctly by the standard Multinomial Naive Bayes algorithm. This implies that most of the positive reviews are classified correctly by the model.

The model predicts lots of reviews as positive when they are actually negative. This results in large FP (False Positive = 2360)

- **(Question 2 - 18 Points)** One potential issue with the original Multinomial Naive Bayes formulation is that if a test instance contains a word that was not seen in training, then $\Pr(word|label) = 0$. To mitigate this issue, one can employ *Laplace Smoothing*. As discussed in class, this implies replacing the standard way of estimating the probability of a word w_k , given a class y_i , with the following equation:

$$\Pr(w_k | y_i) = \frac{n(w_k, y_i) + 1}{\sum_{s=1}^{|V|} n(w_s, y_i) + |V|}. \quad (4)$$

More generally, Laplace Smoothing can be performed according to a parametric equation, where instead of adding 1 to the numerator, we adjust the probability of a word belonging to a class by adding a user-defined parameter α to the numerator, as follows:

$$\Pr(w_k | y_i) = \frac{n(w_k, y_i) + \alpha}{\sum_{s=1}^{|V|} n(w_s, y_i) + \alpha|V|}. \quad (5)$$

Intuitively, setting $\alpha = 0$ results in the standard formulation of Naive Bayes—which does not tackle the problem of words that do not appear in the training set. Suppose, alternatively, that we set $\alpha = 4$. This is equivalent to adding 4 “fake” occurrences of that word to the training set, in order to avoid the zero-frequency problem. Using $\alpha = 1000$, on the other hand, is equivalent to pretending we have seen that word 1000 times in the training set—even though we may have seen it, say, only 8 times. Although this solves the problem of zero-frequency words, it also strongly biases the model to “believe” that that word appears much more frequently than it actually does; and this could make the predictions made by the system less accurate. For these reasons, although it is important/necessary to perform Laplace Smoothing, we have to carefully pick the value of α that works best for our dataset. Using $\alpha = 1$ is common, but other values might result in better performance, depending on the dataset being analyzed.

Another potential issue with the original Multinomial Naive Bayes formulation is that estimating the posterior probabilities using Eq. (1) can lead to numerical instability. This occurs, as discussed in class, because the computation involves multiplying hundreds or thousands of small terms, causing the probability estimate to rapidly approach zero. To tackle this issue, we often compare not $\Pr(y_1 | Doc)$ and $\Pr(y_2 | Doc)$, but the corresponding log-probabilities: $\log(\Pr(y_1 | Doc))$ and $\log(\Pr(y_2 | Doc))$. Taking the logarithm of such probabilities transforms the product of hundreds of terms into the sum of hundreds of terms—which avoids numerical issues. Importantly, it does not change which class is more likely according to the trained model. Concretely, when classifying a new instance, the log-probabilities that should be compared, for each class y_i , are as follows:

$$\log(\Pr(y_i | Doc)) = \log \left(\Pr(y_i) \prod_{k=1}^{len(Doc)} \Pr(w_k | y_i) \right) \quad (6)$$

$$= \log(\Pr(y_i)) + \sum_{k=1}^{len(Doc)} \log(\Pr(w_k | y_i)). \quad (7)$$

In this experiment, you should use 20% of the training set and 20% of the test set; i.e., call the dataset-loading functions by passing 0.2 as their parameters. You should classify the instances in the test set by comparing log-probabilities, $\log(\Pr(y_i | Doc))$, according to Eq. (7), for both classes. You should first report the confusion matrix, precision, recall, and accuracy of your classifier (when evaluated on the test set) when using $\alpha = 1$. Now, vary the value of α from 0.0001 to 1000, by multiplying α by 10 each time. That is, try values of α equal to 0.0001, 0.001, 0.01, 0.1, 1, 10, 100, and 1000. For each value, record the accuracy of the resulting model when evaluated on the test set. Then, create a plot of the model’s accuracy on the test set (shown on the y-axis) as a function of the value of α (shown on the x-axis). The x-axis should represent α values and use a log scale. Analyze this graph and discuss why you think the accuracy suffers when α is too high or too low.

Confusion Matrix resulting from Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities for $\alpha = 1$:

True Positive (TP): 2508

False Negative (FN): 507

False Positive (FP): 334

True Negative (TN): 2641

Accuracy of the model: 0.859599

Precision: 0.882477

Recall: 0.831841

Accuracy of the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities when evaluated on test set for each value of α

Value of Alpha = 0.0001: Accuracy of the Model = 0.798497

Value of Alpha = 0.001: Accuracy of the Model = 0.809182

Value of Alpha = 0.01: Accuracy of the Model = 0.829382
 Value of Alpha = 0.1: Accuracy of the Model = 0.846745
 Value of Alpha = 1: Accuracy of the Model = 0.859599
 Value of Alpha = 10: Accuracy of the Model = 0.859933
 Value of Alpha = 100: Accuracy of the Model = 0.838898
 Value of Alpha = 1000: Accuracy of the Model = 0.813523




Figure 2: Accuracy vs Alpha for Multinomial Naive Bayes

There are 3015 instances of positive reviews. Out of it, 2508 were classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities

There are 2975 instances of negative reviews. Out of it, 2641 were classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities

85.9599% of the reviews are correctly classified by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities

88.32477% of the positive predictions by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities are correct

83.1841% of the positive reviews are classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities

As shown in the graph, the accuracy of the model is highest (85.9933%) when the value of alpha is 10. Other values of alpha strongly biases the model to “believe” that a word appears much more frequently (when $\alpha \geq 100$) or the word appears less frequently (when $\alpha \leq 0.1$) than it actually does. This makes the predictions made by the system less accurate. When $\alpha \leq 0.1$, smoothing is insufficient and when $\alpha \geq 100$, oversmoothing causes the probabilities of the missing words to dominate and bias the model

- **(Question 3 - 18 Points)** Now you will investigate the impact of the training set size on the performance of the model. The classification of new instances, here, should be done by comparing the posterior log-probabilities, $\log(\Pr(y_i | Doc))$, according to Eq. (7), for both classes. You should use the value of α that resulted in the highest accuracy according to your experiments in the previous question. In this question, you should use 100% of the training set and 100% of the test set; i.e., call the dataset-loading

functions by passing 1.0 as their parameters. Then, report (i) the accuracy of your model; (ii) its precision; (iii) its recall; and (iv) the confusion matrix resulting from this experiment.

Impact of the training set size on the performance of the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities

Confusion Matrix resulting from the experiment with 100% of Training Set and 100% of Test Set

True Positive (TP): 12340

False Negative (FN): 2660

False Positive (FP): 1345

True Negative (TN): 13655

Results of the Model with 100% of Training Set and 100% of Test Set

Accuracy of the model: 0.866500

Precision: 0.901717

Recall: 0.822667

There are 15000 instances of positive reviews. Out of it, 12340 were classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 100% of Training Set and 100% of Test Set

There are 15000 instances of negative reviews. Out of it, 13655 were classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 100% of Training Set and 100% of Test Set

86.65% of the reviews are correctly classified by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 100% of Training Set and 100% of Test Set

90.1717% of the positive predictions by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities are correct with 100% of Training Set and 100% of Test Set

82.2667% of the positive reviews are classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 100% of Training Set and 100% of Test Set

The model achieves higher accuracy and precision with the full training dataset

- **(Question 4 - 18 Points)** Now repeat the experiment above but use only 30% of the training instances; that is, load the training set by calling `load_training_set(0.3, 0.3)`. *The entire test set should be used.* Report the same quantities as in the previous question. Discuss whether using such a smaller training set had any impact on the performance of your learned model. Analyze the confusion matrices (of this question and the previous one) and discuss whether one particular class was more affected by changing the size of the training set.

Impact of the smaller training set size (30%) on the performance of the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities

Confusion Matrix resulting from the experiment with 30% of Training Set and 100% of Test Set

True Positive (TP): 11847

False Negative (FN): 3153

False Positive (FP): 1210

True Negative (TN): 13790

Results of the Model with 30% of Training Set and 100% of Test Set Accuracy of the model: 0.854567

Precision: 0.907329

Recall: 0.789800

There are 15000 instances of positive reviews. Out of it, 11847 were classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 30% of Training Set and 100% of Test Set

There are 15000 instances of negative reviews. Out of it, 13790 were classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 30% of Training Set and 100% of Test Set

85.4567% of the reviews are correctly classified by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 30% of Training Set and 100% of Test Set

90.7329% of the positive predictions by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities are correct with 30% of Training Set and 100% of Test Set

78.9800% of the positive reviews are classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 30% of Training Set and 100% of Test Set

When the training set size is reduced from 100% to 30%, the following changes are observed in the confusion matrix

The True positive value has fallen from 12340 to 11847

The False Negative value has risen from 2660 to 3153

The predictions of the positive class are impacted more by reducing the training set size from 100% to 30%

The False positive value has fallen from 1345 to 1210

The True Negative value has risen from 13655 to 13790

The predictions of the negative class are less impacted by reducing the training set size from 100% to 30%

Also the value of Recall has dropped from 0.822667 to 0.789800. This implies that the percentage of correct classification by the model for the positive class has reduced. This implies the model has become less effective.

- **(Question 5 - 10 Points)** In the document-classification application studied in this homework (i.e., one where ML is used to accurately classify product reviews), would you say that it is more important to have high accuracy, high precision, or high recall? Justify your opinion.

In this example, I would say that it is more important to have high recall than high precision or high accuracy. Recall tells us, out of all positive reviews, how many the Multinomial model is able to correctly identify as positive. This percentage represents the correct classification of positive reviews by the model. High Recall helps ensure that most positive reviews are detected. If the model can correctly identify the positive reviews, it will help the organization to better understand customer satisfaction and market those products more effectively

- **(Question 6 - 18 Points)** Finally, you will study how the performance of the learned model is affected by training it using an *unbalanced* dataset (i.e., a dataset with significantly more examples of one of the classes). The classification of new instances, here, should be done by comparing the posterior log-probabilities, $\log(\Pr(y_i | Doc))$, according to Eq. (7), for both classes. You should use the value of α that resulted in the highest accuracy according to your experiments in the previous questions. You will now conduct an experiment where you use only 10% of the available *positive* training instances and that uses 50% of the available *negative* training instances. That is, use `load_training_set(0.1, 0.5)`. *The entire test set should be used.* Show the confusion matrix of your trained model, as well as its accuracy, precision, and recall. Compare this model's performance to the performance (according to these same metrics) of the model trained in question Q.4—that is, a model that was trained with a *balanced* dataset. Discuss how training under an unbalanced dataset affected each of these performance metrics.

Impact of the unbalanced training dataset on the performance of the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities

Confusion Matrix resulting from the experiment with Unbalanced Training Set (10% of available positive and 50% of available negative training instances) and 100% of Test Set

True Positive (TP): 44
False Negative (FN): 14956
False Positive (FP): 28
True Negative (TN): 14972

Results of the Model with Unbalanced Training Set (10% of available positive and 50% of available negative training instances) and 100% of Test Set

Accuracy of the model: 0.500533
Precision: 0.611111
Recall: 0.002933

There are 15000 instances of positive reviews. Out of it, 44 were classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 10% of Positive Training Set and 50% of Negative Training set and 100% of Test Set

There are 15000 instances of negative reviews. Out of it, 14972 were classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 10% of Positive Training Set and 50% of Negative Training set and 100% of Test Set

50.0533% of the reviews are correctly classified by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities with 10% of Positive Training Set and 50% of Negative Training set and 100% of Test Set

61.1111% of the positive predictions by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities are correct with 10% of Positive Training Set and 50% of Negative Training set and 100% of Test Set

00.2933% of the positive reviews are classified correctly by the Multinomial Naive Bayes algorithm with Laplace Smoothing and Log-Probabilities 10% of Positive Training Set and 50% of Negative Training set and 100% of Test Set

When the training set size is changed to Unbalanced dataset with 10% of Positive Training Set and 50% of Negative Training set, the following changes are observed in the confusion matrix

True positive value has fallen from 11847 to 44

False negative value has risen from 3153 to 14956

The value of Recall has dropped from 0.789800 to 0.002933. This implies the model has become less effective in predicting the positive reviews as positive

False positive value has fallen from 1210 to 28

The value of Precision has dropped from 0.907329 to 0.611111. It looks like the model has lost the ability to make positive predictions. It could be due to using only 10% of the training set as positive reviews.

True Negative has risen from 13790 to 14972

It looks like the model incorrectly predicts all instances as negative class, rarely identifying positive classes

Because the model is trained with far fewer positive reviews, the model has deteriorated significantly in its ability to identify positive reviews and identifies most reviews as negative reviews